

***** x86 ASSEMBLY BASICS *****

```
0 "TOPICS OF DISCUSSION"  
1 "WHAT IS ASSEMBLY"  
2 "REGISTERS"  
3 "INSTRUCTIONS"  
4 "MAKING A+B PROGRAM"  
5 "PASSING IT ARGUMENTS"  
6 "SYSTEM CALLS"  
7 "MAKING PUT-STRING"  
8 "THE STACK"  
9 "DISASSEMBLING C"  
10 "HARDWARE INTERACTION"
```



The slides & practice programs can be found on github

github.com/susikohmelo/assembly_basics_presentation



What **actually** is Assembly?

Generic term

Languages which communicate directly with hardware.

Hardware specific.

Many different syntax can exist for the same language

***** EXAMPLES *****

x86 : intel & AMD

ARM : mobile

RISCV : open-source
alternative

(Almost) everything is just memory

Operations, variables...
It's all just data in memory.

Fundamentally there is no difference between data and code.

Everything in your program is just a value stored in some address of RAM which you can freely access and edit.
(Assuming the OS lets you)

You can even make self modifying code



Registers

Super fast, temporary memory inside the CPU

Registers can be used to:

- Store/assign values from/to other registers & RAM
- Do operations on the value they hold (add, multiply etc.)

(Some) commonly used registers

***** GENERAL *****

ANYTHING GOES HERE

ax, bx, cx, dx

SRC/DEST POINTERS

si : source index

di : destination

***** SPECIAL *****

STACK POINTERS

sp : Top pointer

bp : Bottom pointer

FLAG REGISTERS

bazillion of them

Register inception

Registers can be divided into multiple smaller registers



Instructions

Instructions are the actual commands we can give the CPU to execute

- Operations on registers (assign, add, multiply etc.)
- Move values in RAM & registers
- Jump to locations

***** EXAMPLES *****

MOV **ax, 42**

Move the value 42
into register ax

ADD **ax, bx**

Add value in bx to
the value in ax.

Simple program example

***** PROGRAM *****

_start: 

MOV ax, 42

MOV bx, ax

ADD ax, bx

MOV bx, 0x42

MOV [bx], ax

***** REGISTERS *****

ax : 0

bx : 0

***** RAM *****

0x42 = 0

Simple program example

***** PROGRAM *****

_start:

MOV ax, 42 ←

MOV bx, ax

ADD ax, bx

MOV bx, 0x42

MOV [bx], ax

***** REGISTERS *****

ax : 42

bx : 0

***** RAM *****

0x42 = 0

Simple program example

***** PROGRAM *****

_start:

MOV ax, 42

MOV bx, ax ←

ADD ax, bx

MOV bx, 0x42

MOV [bx], ax

***** REGISTERS *****

ax : 42

bx : 42

***** RAM *****

0x42 = 0

Simple program example

***** PROGRAM *****

_start:

MOV ax, 42

MOV bx, ax

ADD ax, bx ←

MOV bx, 0x42

MOV [bx], ax

***** REGISTERS *****

ax : 84

bx : 42

***** RAM *****

0x42 = 0

Simple program example

***** PROGRAM *****

_start:

MOV ax, 42

MOV bx, ax

ADD ax, bx

MOV bx, 0x42 ←

MOV [bx], ax

***** REGISTERS *****

ax : 84

bx : 0x42

***** RAM *****

0x42 = 0

Simple program example

***** PROGRAM *****

_start:

MOV ax, 42

MOV bx, ax

ADD ax, bx

MOV bx, 0x42

MOV [bx], ax ←

***** REGISTERS *****

ax : 84

bx : 42

***** RAM *****

0x42 = 84

**Let's make a simple A+B function that
can be called from a C program**

How can we pass parameters/args?

Many ways!

1. Arbitrary RAM location (bad)
2. Registers
3. Stack

In 32-bit C programs, all parameters are pushed onto stack.
[We'll look at this later]

In 64-bit C programs, parameters are put into specific registers, arguments > 6 get put onto stack.

**** 64-BIT LINUX **** **CALLING CONVENTION**

rdi	:	param.	1
rsi	:	param.	2
rdx	:	param.	3
rcx	:	param.	4
r8	:	param.	5
r9	:	param.	6
stack	:	param.	7..

Let's add that

System calls (syscall)

System calls ask the OS to do something for us

Usually done by sending a software interrupt signal.

Registers are used as parameters for the syscall.

‘syscall’ sends the interrupt

```
**** EXAMPLE ****  
  
/* syscall #60 */  
MOV     rax, 60  
  
/* exit code */  
MOV     rdi, 42  
  
syscall
```

"Simple" register-register op (ADD,OR,etc.)

<1

Memory write

~1

Bypass delay: switch between
integer and floating-point units

0-3

"Right" branch of "if"

1-2

Floating-point/vector addition

1-3

Multiplication (integer/float/vector)

1-7

Return error and check

1-7

L1 read

3-4

TLB miss

7-21

L2 read

10-12

"Wrong" branch of "if" (branch misprediction)

10-20

Floating-point division

10-40

128-bit vector division

10-70

Atomics/CAS

15-30

C function direct call

15-30

Integer division

15-40

C function indirect call

20-50

C++ virtual function call

30-60

L3 read

30-70

Main RAM read

100-150

NUMA: different-socket atomics/CAS
(guesstimate)

100-300

NUMA: different-socket L3 read

100-300

Allocation+deallocation pair (small objects)

200-500

NUMA: different-socket main RAM read

300-500

Kernel call

1000-1500

Thread context switch (direct costs)

2000

C++ Exception thrown+caught

5000-10000

Thread context switch (total costs,
including cache invalidation)

10000 - 1 million

System calls are expensive



**Let's print something onto the screen
with a system call**

The stack redpill

sp/bp point to the top and bottom of the stack

You can access anything in the stack, not just the top!

When you push a value, it gets put in to the address that the stack top is pointing to & the pointer is incremented.

The stack grows downwards. So as you add more things, the top of the stack becomes a smaller address.

**You can PUSH and POP registers and values onto the stack.
But you can also just use the stack like a normal pointer!**

NOTE

In the examples the push/pop operations add/remove 1 byte.

**Unless you're on 8/16-bit hardware,
the stack typically operates on 4+ bytes.**

The return address is also not being shown
[That is coming later]

Simple program example

***** PROGRAM *****

_start: 


PUSH 42
PUSH 3

MOV a1, [sp+1]
POP b1

ADD sp, 1

***** STACK *****

RAM

0x00	:	0	
0x01	:	0	
0x02	:	0	
0x03	:	0	 sp/bp

REGISTERS

sp	:	0x03
bp	:	0x03
a1	:	0
b1	:	0

Simple program example

***** PROGRAM *****

_start:

PUSH 42 ←
PUSH 3

MOV a1, [sp+1]
POP b1

ADD sp, 1

***** STACK *****

RAM

0x00 : 0
0x01 : 0
0x02 : 42 ← sp
0x03 : 0 ← bp

REGISTERS

sp : 0x02
bp : 0x03
a1 : 0
b1 : 0

Simple program example

***** PROGRAM *****

_start:

PUSH 42

PUSH 3 ←

MOV a1, [sp+1]

POP b1

ADD sp, 1

***** STACK *****

RAM

0x00 : 0

0x01 : 3 ← sp

0x02 : 42

0x03 : 0 ← bp

REGISTERS

sp : 0x01

bp : 0x03

a1 : 0

b1 : 0

Simple program example

**** PROGRAM ****

_start:

PUSH 42
PUSH 3

MOV a1, [sp+1]
POP b1

ADD sp, 1

**** STACK ****

RAM

0x00	:	0	
0x01	:	3	← sp
0x02	:	42	
0x03	:	0	← bp

REGISTERS

sp	:	0x01
bp	:	0x03
a1	:	42
b1	:	0

Simple program example

**** PROGRAM ****

_start:

PUSH 42
PUSH 3

MOV a1, [sp+1]
POP b1 ←

ADD sp, 1

**** STACK ****

RAM

0x00 : 0
0x01 : 3
0x02 : 42 ← sp
0x03 : 0 ← bp

REGISTERS

sp : 0x02
bp : 0x03
a1 : 42
b1 : 3

Simple program example

***** PROGRAM *****

_start:

PUSH 42

PUSH 3

MOV a1, [sp+1]

POP b1

ADD sp, 1 ←

***** STACK *****

RAM

0x00 : 0

0x01 : 3

0x02 : 42

0x03 : 0 ← sp/bp

REGISTERS

sp : 0x03

bp : 0x03

a1 : 42

b1 : 3

‘Stack frames’

Refers to a region of the stack “belonging” to a function

On function calls, the bottom ptr is saved to the stack.
The bottom ptr is then set to point to top of the stack.

This effectively makes a new stack on top of the old one.

When the function returns, we pop the old bottom pointers location and now we are back in the original stack frame.

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...

    POP     bp
    RET

_start:                                     ←
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0
0x01	:	0
0x02	:	0
0x03	:	0
0x04	:	0
0x05	:	0
		← sp/bp

REGISTERS

sp	:	0x05
bp	:	0x05

Simple program example

***** PROGRAM *****

```
my_funct:  
    PUSH    bp  
    MOV     bp, sp  
    ...  
  
    POP     bp  
    RET  
  
_start:  
    PUSH    42  
    CALL    my_funct()  
    ...
```

***** STACK *****

RAM

0x00	:	0
0x01	:	0
0x02	:	0
0x03	:	0
0x04	:	42
0x05	:	0

↑ sp
↑ bp

REGISTERS

sp	:	0x04
bp	:	0x05

Simple program example

***** PROGRAM *****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...

    POP     bp
    RET

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

***** STACK *****

RAM

0x00	:	0	
0x01	:	0	
0x02	:	0	
0x03	:	0	
0x04	:	42	← sp
0x05	:	0	← bp

REGISTERS

sp	:	0x04
bp	:	0x05

Simple program example

***** PROGRAM *****

```
my_func↑:  
    PUSH    bp  
    MOV     bp, sp  
    ...  
  
    POP     bp  
    RET  
  
_start:  
    PUSH    42  
    CALL    my_func()  
    ...
```

***** STACK *****

RAM

0x00	:	0	
0x01	:	0	
0x02	:	0	
0x03	:	0	
0x04	:	42	← sp
0x05	:	0	← bp

REGISTERS

sp	: 0x04
bp	: 0x05

Simple program example

**** PROGRAM ****

```
my_funct:  
  PUSH bp      ←  
  MOV  bp, sp  
  ...  
  POP  bp  
  RET  
  
_start:  
  PUSH 42  
  CALL my_funct()  
  ...
```

**** STACK ****

RAM

0x00	:	0	
0x01	:	0	
0x02	:	0	
0x03	:	0x05	← sp
0x04	:	42	
0x05	:	0	← bp

REGISTERS

sp	:	0x03
bp	:	0x05

Simple program example

***** PROGRAM *****

```
my_funct:  
    PUSH    bp  
    MOV     bp, sp    ←  
    ...  
    POP     bp  
    RET  
  
_start:  
    PUSH    42  
    CALL    my_funct()  
    ...
```

***** STACK *****

RAM

0x00	:	0
0x01	:	0
0x02	:	0
0x03	:	0x05 ← sp/bp
0x04	:	42
0x05	:	0

REGISTERS

sp	:	0x03
bp	:	0x03

Simple program example

***** PROGRAM *****

```
my_func+t:  
  PUSH  bp  
  MOV   bp, sp  
  ...  
  POP   bp  
  RET  
  
_start:  
  PUSH  42  
  CALL  my_func+t()  
  ...
```

***** STACK *****

RAM

0x00	:	0
0x01	:	0
0x02	:	0
0x03	:	0x05 ⁺ sp/bp
0x04	:	42
0x05	:	0

REGISTERS

sp	:	0x03
bp	:	0x03

Simple program example

***** PROGRAM *****

```
my_func+t:
    PUSH    bp
    MOV     bp, sp
    ...
    POP     bp
    RET

_start:
    PUSH    42
    CALL    my_func+t()
    ...
```

***** STACK *****

RAM

0x00	:	0	
0x01	:	0	
0x02	:	99	⁺ sp
0x03	:	0x05	⁺ bp
0x04	:	42	
0x05	:	0	

REGISTERS

sp	:	0x02
bp	:	0x03

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...
    POP     bp
    RET

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0	
0x01	:	22	+ sp
0x02	:	99	
0x03	:	0x05	+ bp
0x04	:	42	
0x05	:	0	

REGISTERS

sp	:	0x01
bp	:	0x03

Simple program example

**** PROGRAM ****

```
my_func+t:  
    PUSH    bp  
    MOV     bp, sp  
    ...  
    POP     bp  
    RET  
  
_start:  
    PUSH    42  
    CALL    my_func+t()  
    ...
```

**** STACK ****

RAM

0x00	:	0	
0x01	:	22	
0x02	:	99	⁺ sp
0x03	:	0x05	⁺ bp
0x04	:	42	
0x05	:	0	

REGISTERS

sp	:	0x02
bp	:	0x03

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...
    POP     bp
    RET

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0
0x01	:	22
0x02	:	99
0x03	:	0x05+ sp/bp
0x04	:	42
0x05	:	0

REGISTERS

sp	:	0x03
bp	:	0x03

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...

    POP     bp          ←
    RET

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0
0x01	:	22
0x02	:	99
0x03	:	0x05
0x04	:	42 ← sp
0x05	:	0 ← bp

REGISTERS

sp	:	0x04
bp	:	0x05

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...

    POP     bp
    RET                               ←

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0
0x01	:	22
0x02	:	99
0x03	:	0x05
0x04	:	42 ← sp
0x05	:	0 ← bp

REGISTERS

sp	:	0x04
bp	:	0x05

Simple program example

**** PROGRAM ****

```
my_func_t:
    PUSH    bp
    MOV     bp, sp
    ...

    POP     bp
    RET

_start:
    PUSH    42
    CALL    my_func_t()
    ...
```

**** STACK ****

RAM

0x00	:	0	
0x01	:	22	
0x02	:	99	
0x03	:	0x05	
0x04	:	42	← sp
0x05	:	0	← bp

REGISTERS

sp	:	0x04
bp	:	0x05

Let's look at some disassembled C code

**We should now be able to understand
quite a lot of what's going on.**

How does the CPU know where to go

Push, then pop a return address

A pointer to the current line of code being executed by the CPU is kept in the instruction pointer register (IP)

When you call a function:

1. Pushes next line of code to the stack
2. Sets IP as the newly called address

Returning from a function

1. Pops the address back into IP



Interacting with hardware

The majority of devices are mapped somewhere in memory

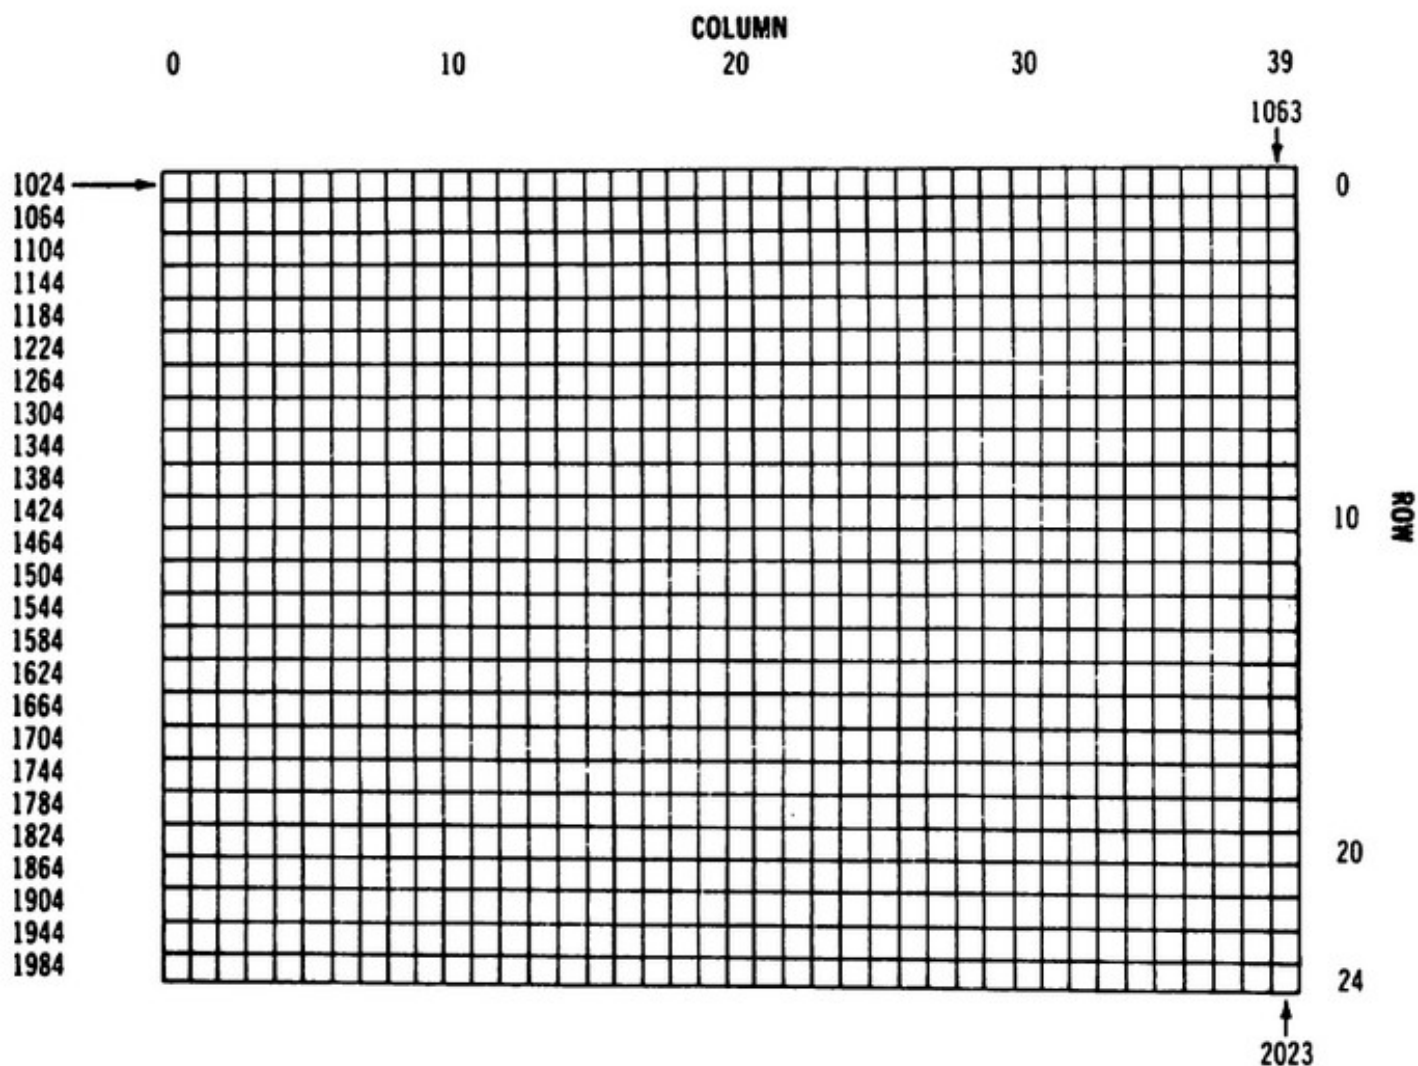
You can communicate with devices by reading/writing some specific area of memory, usually in RAM.

In addition to RAM, there is I/O memory.

You can think of it like a completely separate mini-RAM. The way you use it is the same as RAM. But instead of the 'mov' instruction, you use 'in' and 'out' to read/write.

I/O memory is often used to configure the hardware. Like telling the device which RAM address it should map to.

SCREEN MEMORY MAP



Map of the
screen from C64

To put a char on
screen, move the
value to the
corresponding
location in RAM.

Example:

```
MOV edi, 1024  
MOV [edi], 13
```

Would display
the character
corresponding to
the value 13 at
the top left of
the screen

****Though if you
were actually
using a C64, you
would use BASIC
command 'POKE'.**

The slides & practice programs can be found on github

github.com/susikohmelo/assembly_basics_presentation

